

ai-ml-tasks

June 5, 2025

https://colab.research.google.com/drive/1SR0yLmwUS68UMHsDGHSB1Ofzt_kNPzYY?usp=sharing

1 Github Repository

```
[61]: %cd https://github.com/zoya4477/AI-Ml.git
!git clone
!git config --global user.email "zoyahafeez785@gmail.com"
!git config --global user.name "zoya4477"
```

```
[Errno 2] No such file or directory: 'https://github.com/zoya4477/AI-Ml.git'
/content
```

```
fatal: You must specify a repository to clone.
```

```
usage: git clone [<options>] [--] <repo> [<dir>]
```

-v, --verbose	be more verbose
-q, --quiet	be more quiet
--progress	force progress reporting
--reject-shallow	don't clone shallow repository
-n, --no-checkout	don't create a checkout
--bare	create a bare repository
--mirror	create a mirror repository (implies bare)
-l, --local	to clone from a local repository
--no-hardlinks	don't use local hardlinks, always copy
-s, --shared	setup as shared repository
--recurse-submodules[=<pathspec>]	initialize submodules in the clone
--recursive ...	alias of --recurse-submodules
-j, --jobs <n>	number of submodules cloned in parallel
--template <template-directory>	directory from which templates will be used
--reference <repo>	reference repository
--reference-if-able <repo>	reference repository
--dissociate	use --reference only while cloning
-o, --origin <name>	use <name> instead of 'origin' to track upstream
-b, --branch <branch>	

```

        checkout <branch> instead of the remote's HEAD
-u, --upload-pack <path>
                        path to git-upload-pack on the remote
--depth <depth>        create a shallow clone of that depth
--shallow-since <time>
                        create a shallow clone since a specific time
--shallow-exclude <revision>
                        deepen history of shallow clone, excluding rev
--single-branch         clone only one branch, HEAD or --branch
--no-tags               don't clone any tags, and make later fetches not to
follow them
--shallow-submodules    any cloned submodules will be shallow
--separate-git-dir <gitdir>
                        separate git dir from working tree
-c, --config <key=value>
                        set config inside the new repository
--server-option <server-specific>
                        option to transmit
-4, --ipv4              use IPv4 addresses only
-6, --ipv6              use IPv6 addresses only
--filter <args>         object filtering
--remote-submodules     any cloned submodules will use their remote-tracking
branch
--sparse                initialize sparse-checkout file to include only files
at root

```

2 Task 1: Exploring and Visualizing the Iris Dataset

3 Load the Dataset

```

[30]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Load the Iris dataset directly from seaborn
df = pd.read_csv('/content/Iris.csv')

```

4 Inspect the dataset

```

[31]: # Shape of the dataset
print("Shape of the dataset:", df.shape)

# Column names
print("Column names:", df.columns.tolist())

```

```

# First few rows
print(df.head())

# Info summary
print("\nDataset Info:")
print(df.info())

# Descriptive statistics
print("\nDescriptive Statistics:")
print(df.describe())

```

Shape of the dataset: (152, 6)

Column names: ['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm', 'Species']

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	NaN	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	NaN	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	NaN	Iris-setosa
4	5	5.0	NaN	1.4	0.2	Iris-setosa

Dataset Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 152 entries, 0 to 151

Data columns (total 6 columns):

#	Column	Non-Null Count	Dtype
0	Id	152 non-null	int64
1	SepalLengthCm	151 non-null	float64
2	SepalWidthCm	151 non-null	float64
3	PetalLengthCm	151 non-null	float64
4	PetalWidthCm	151 non-null	float64
5	Species	152 non-null	object

dtypes: float64(4), int64(1), object(1)

memory usage: 7.3+ KB

None


Descriptive Statistics:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
count	152.000000	151.000000	151.000000	151.000000	151.000000
mean	75.414474	5.849007	3.055629	3.770861	1.206623
std	43.866813	0.823073	0.432302	1.764902	0.766870
min	1.000000	4.300000	2.000000	1.000000	0.100000
25%	37.750000	5.100000	2.800000	1.600000	0.300000
50%	75.500000	5.800000	3.000000	4.400000	1.300000
75%	113.250000	6.400000	3.300000	5.100000	1.800000
max	150.000000	7.900000	4.400000	6.900000	2.500000

```
#Visualize the Dataset
```

```
[32]: #Scatter Plot -- Sepal Length vs Sepal Width
```

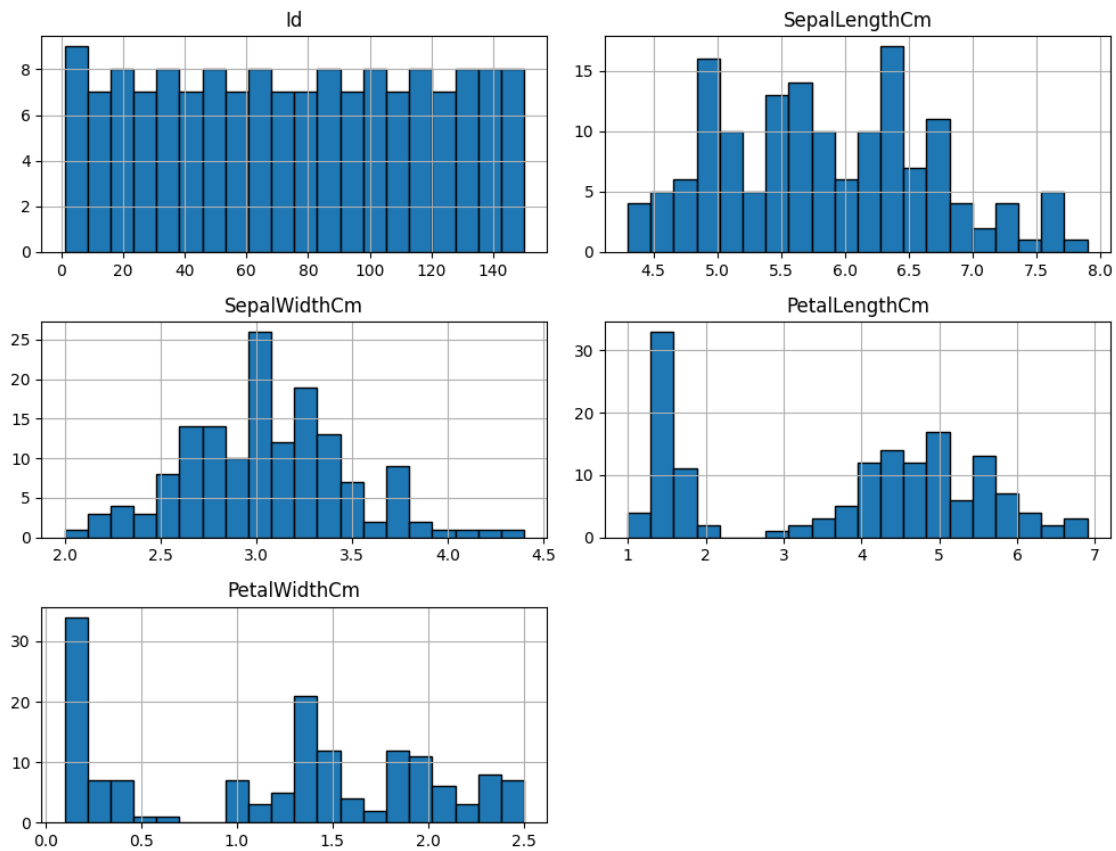
```
plt.figure(figsize=(8, 6))  
sns.scatterplot(data=df, x='SepalLengthCm', y='SepalWidthCm', hue='Species')  
plt.title('Sepal Length vs Sepal Width')  
plt.show()
```



```
[33]: #Histograms -- Distribution of Each Feature
```

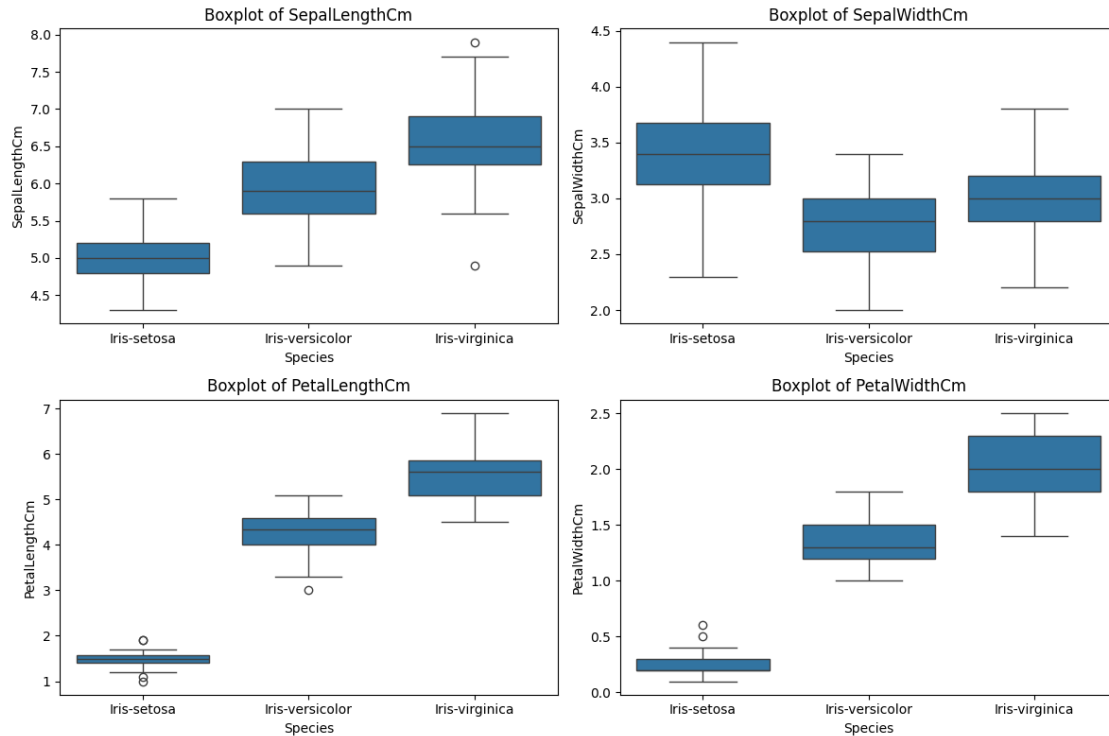
```
df.hist(figsize=(10, 8), bins=20, edgecolor='black')  
plt.suptitle('Feature Distributions')  
plt.tight_layout()  
plt.show()
```

Feature Distributions



```
[34]: #Boxplot -- To Identify Outlier
plt.figure(figsize=(12, 8))
features = ['SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm']

for i, column in enumerate(features):
    plt.subplot(2, 2, i + 1)
    sns.boxplot(x='Species', y=column, data=df)
    plt.title(f'Boxplot of {column}')
plt.tight_layout()
plt.show()
```



[34]:

5 Task 2: Predict Future Stock Prices (Short-Term)

[35]: `pip install yfinance`

```
Requirement already satisfied: yfinance in /usr/local/lib/python3.11/dist-packages (0.2.61)
Requirement already satisfied: pandas>=1.3.0 in /usr/local/lib/python3.11/dist-packages (from yfinance) (2.2.2)
Requirement already satisfied: numpy>=1.16.5 in /usr/local/lib/python3.11/dist-packages (from yfinance) (2.0.2)
Requirement already satisfied: requests>=2.31 in /usr/local/lib/python3.11/dist-packages (from yfinance) (2.32.3)
Requirement already satisfied: multitasking>=0.0.7 in /usr/local/lib/python3.11/dist-packages (from yfinance) (0.0.11)
Requirement already satisfied: platformdirs>=2.0.0 in /usr/local/lib/python3.11/dist-packages (from yfinance) (4.3.8)
Requirement already satisfied: pytz>=2022.5 in /usr/local/lib/python3.11/dist-packages (from yfinance) (2025.2)
Requirement already satisfied: frozendict>=2.3.4 in /usr/local/lib/python3.11/dist-packages (from yfinance) (2.4.6)
Requirement already satisfied: peewee>=3.16.2 in /usr/local/lib/python3.11/dist-packages (from yfinance) (3.16.2)
```

packages (from yfinance) (3.18.1)
 Requirement already satisfied: beautifulsoup4>=4.11.1 in
 /usr/local/lib/python3.11/dist-packages (from yfinance) (4.13.4)
 Requirement already satisfied: curl_cffi>=0.7 in /usr/local/lib/python3.11/dist-
 packages (from yfinance) (0.11.1)
 Requirement already satisfied: protobuf>=3.19.0 in
 /usr/local/lib/python3.11/dist-packages (from yfinance) (5.29.5)
 Requirement already satisfied: websockets>=13.0 in
 /usr/local/lib/python3.11/dist-packages (from yfinance) (15.0.1)
 Requirement already satisfied: soupsieve>1.2 in /usr/local/lib/python3.11/dist-
 packages (from beautifulsoup4>=4.11.1->yfinance) (2.7)
 Requirement already satisfied: typing-extensions>=4.0.0 in
 /usr/local/lib/python3.11/dist-packages (from beautifulsoup4>=4.11.1->yfinance)
 (4.13.2)
 Requirement already satisfied: cffi>=1.12.0 in /usr/local/lib/python3.11/dist-
 packages (from curl_cffi>=0.7->yfinance) (1.17.1)
 Requirement already satisfied: certifi>=2024.2.2 in
 /usr/local/lib/python3.11/dist-packages (from curl_cffi>=0.7->yfinance)
 (2025.4.26)
 Requirement already satisfied: python-dateutil>=2.8.2 in
 /usr/local/lib/python3.11/dist-packages (from pandas>=1.3.0->yfinance)
 (2.9.0.post0)
 Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-
 packages (from pandas>=1.3.0->yfinance) (2025.2)
 Requirement already satisfied: charset-normalizer<4,>=2 in
 /usr/local/lib/python3.11/dist-packages (from requests>=2.31->yfinance) (3.4.2)
 Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-
 packages (from requests>=2.31->yfinance) (3.10)
 Requirement already satisfied: urllib3<3,>=1.21.1 in
 /usr/local/lib/python3.11/dist-packages (from requests>=2.31->yfinance) (2.4.0)
 Requirement already satisfied: pycparser in /usr/local/lib/python3.11/dist-
 packages (from cffi>=1.12.0->curl_cffi>=0.7->yfinance) (2.22)
 Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-
 packages (from python-dateutil>=2.8.2->pandas>=1.3.0->yfinance) (1.17.0)

#Import libraries and load data

```

[36]: import yfinance as yf
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt

# Download historical data for Apple (AAPL)
ticker = 'AAPL'
  
```

```
df = yf.download(ticker, start='2020-01-01', end='2024-01-01')

# Display first rows
print(df.head())
```

[*****100%*****] 1 of 1 completed

Price	Close	High	Low	Open	Volume
Ticker	AAPL	AAPL	AAPL	AAPL	AAPL
Date					
2020-01-02	72.620834	72.681281	71.373211	71.627084	135480400
2020-01-03	71.914795	72.676423	71.689935	71.847095	146322800
2020-01-06	72.487839	72.526526	70.783241	71.034702	118387200
2020-01-07	72.146950	72.753831	71.926922	72.497537	108872000
2020-01-08	73.307510	73.609745	71.849533	71.849533	132079200

6 Prepare features and target

```
[37]: # Shift the Close column up by 1 to represent next day close price
df['Next_Close'] = df['Close'].shift(-1)

# Drop last row with NaN target
df = df[:-1]

# Features and target
features = ['Open', 'High', 'Low', 'Volume']
X = df[features]
y = df['Next_Close']
```

#Split data into train/test sets

```
[38]: X_train, X_test, y_train, y_test = train_test_split(X, y, shuffle=False,
↳ test_size=0.2)
```

7 Train the model

```
[39]: #Linear Regression
model = LinearRegression()
model.fit(X_train, y_train)
```


```
[39]: LinearRegression()
```

```
[40]: y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error: {mse:.4f}")
```


Mean Squared Error: 4.9761

#Plot actual vs predicted closing prices

```
[41]: plt.figure(figsize=(12, 6))
plt.plot(y_test.index, y_test, label='Actual Close Price', color='blue')
plt.plot(y_test.index, y_pred, label='Predicted Close Price', color='red',
         linestyle='--')
plt.title(f'{ticker} Actual vs Predicted Close Price')
plt.xlabel('Date')
plt.ylabel('Price ($)')
plt.legend()
plt.show()
```



```
[41]:
```

8 Task 3: Heart Disease Prediction

9 Import Libraries

```
[42]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.metrics import accuracy_score, roc_auc_score, roc_curve, \
    confusion_matrix, ConfusionMatrixDisplay
```

10 Load Dataset

```
[43]: data = pd.read_csv('/content/archive.zip')
      data.head()
```

```
[43]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	69	1	0	160	234	1	2	131	0	0.1	1	
1	69	0	0	140	239	0	0	151	0	1.8	0	
2	66	0	0	150	226	0	0	114	0	2.6	2	
3	65	1	0	138	282	1	2	174	0	1.4	1	
4	64	1	0	110	211	0	2	144	1	1.8	1	

	ca	thal	condition
0	1	0	0
1	2	0	0
2	0	0	0
3	1	0	1
4	0	0	0

10.1 Check Missing Values

```
[44]: print("Missing values in each column:")
      print(data.isnull().sum())
```

```
Missing values in each column:
age          0
sex          0
cp           0
trestbps     0
chol         0
fbs          0
restecg      0
thalach      0
exang        0
oldpeak      0
slope        0
ca           0
thal         0
condition    0
dtype: int64
```

11 Basic Info and Description

```
[45]: print(data.info())
      print(data.describe())
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 297 entries, 0 to 296
```

```
Data columns (total 14 columns):
```

#	Column	Non-Null Count	Dtype
0	age	297 non-null	int64
1	sex	297 non-null	int64
2	cp	297 non-null	int64
3	trestbps	297 non-null	int64
4	chol	297 non-null	int64
5	fbs	297 non-null	int64
6	restecg	297 non-null	int64
7	thalach	297 non-null	int64
8	exang	297 non-null	int64
9	oldpeak	297 non-null	float64
10	slope	297 non-null	int64
11	ca	297 non-null	int64
12	thal	297 non-null	int64
13	condition	297 non-null	int64

```
dtypes: float64(1), int64(13)
```

```
memory usage: 32.6 KB
```

```
None
```


	age	sex	cp	trestbps	chol	fbs	\
count	297.000000	297.000000	297.000000	297.000000	297.000000	297.000000	
mean	54.542088	0.676768	2.158249	131.693603	247.350168	0.144781	
std	9.049736	0.468500	0.964859	17.762806	51.997583	0.352474	
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	
25%	48.000000	0.000000	2.000000	120.000000	211.000000	0.000000	
50%	56.000000	1.000000	2.000000	130.000000	243.000000	0.000000	
75%	61.000000	1.000000	3.000000	140.000000	276.000000	0.000000	
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	

	restecg	thalach	exang	oldpeak	slope	ca	\
count	297.000000	297.000000	297.000000	297.000000	297.000000	297.000000	
mean	0.996633	149.599327	0.326599	1.055556	0.602694	0.676768	
std	0.994914	22.941562	0.469761	1.166123	0.618187	0.938965	
min	0.000000	71.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	133.000000	0.000000	0.000000	0.000000	0.000000	
50%	1.000000	153.000000	0.000000	0.800000	1.000000	0.000000	
75%	2.000000	166.000000	1.000000	1.600000	1.000000	1.000000	
max	2.000000	202.000000	1.000000	6.200000	2.000000	3.000000	

	thal	condition
count	297.000000	297.000000
mean	0.835017	0.461279
std	0.956690	0.499340
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	2.000000	1.000000
max	2.000000	1.000000

12 Visualize Target Distribution


```
[48]: sns.countplot(x='condition', data=data)  
plt.title('Target Class Distribution (Heart Disease Presence)')  
plt.show()
```







13 Distribution of Numerical Features

```
[59]: numerical_cols = ['age', 'trestbps', 'chol', 'thalach', 'oldpeak']

for col in numerical_cols:
    plt.figure(figsize=(6, 4))
    sns.histplot(data[col], kde=True, bins=30)
    plt.title(f'Distribution of {col}')
    plt.show()
```



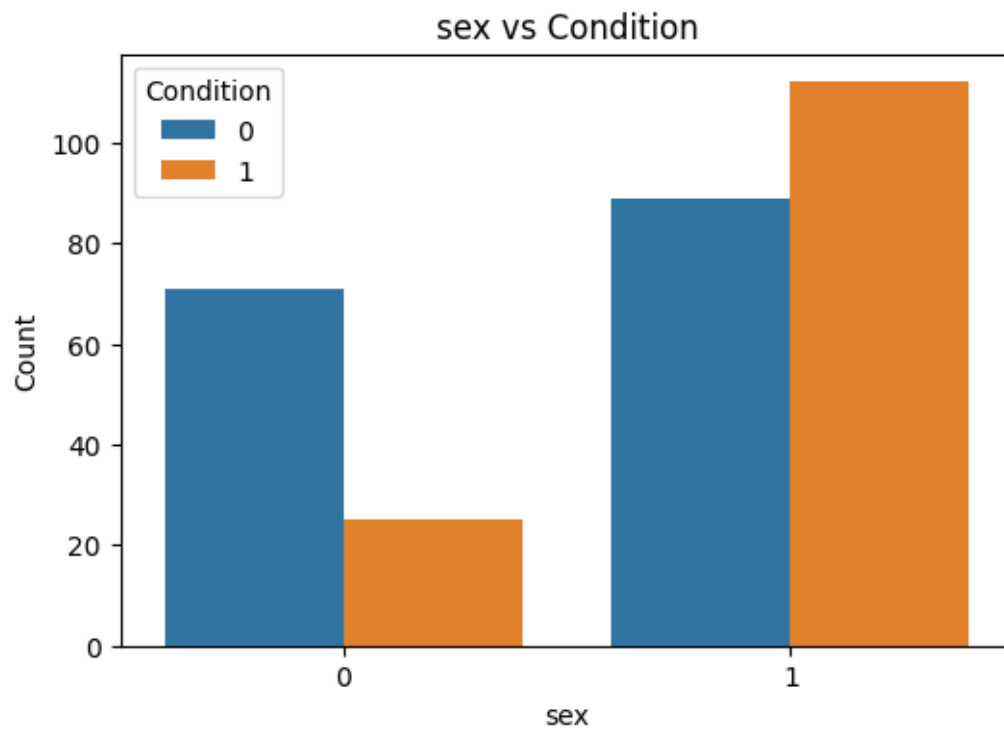







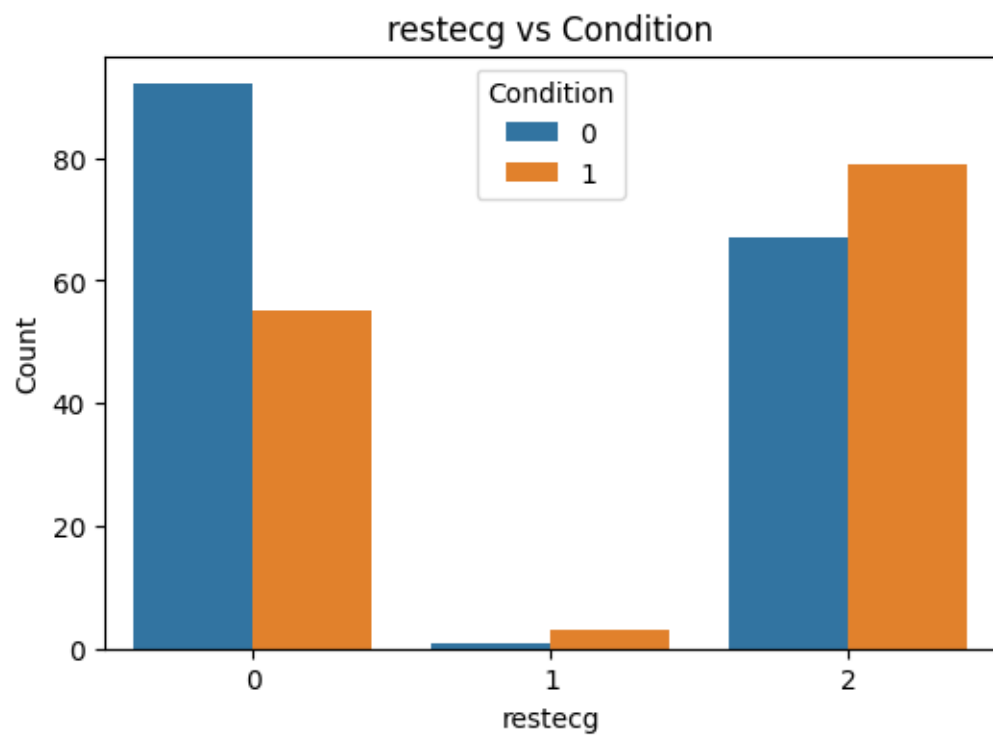
14 Categorical Features vs Target

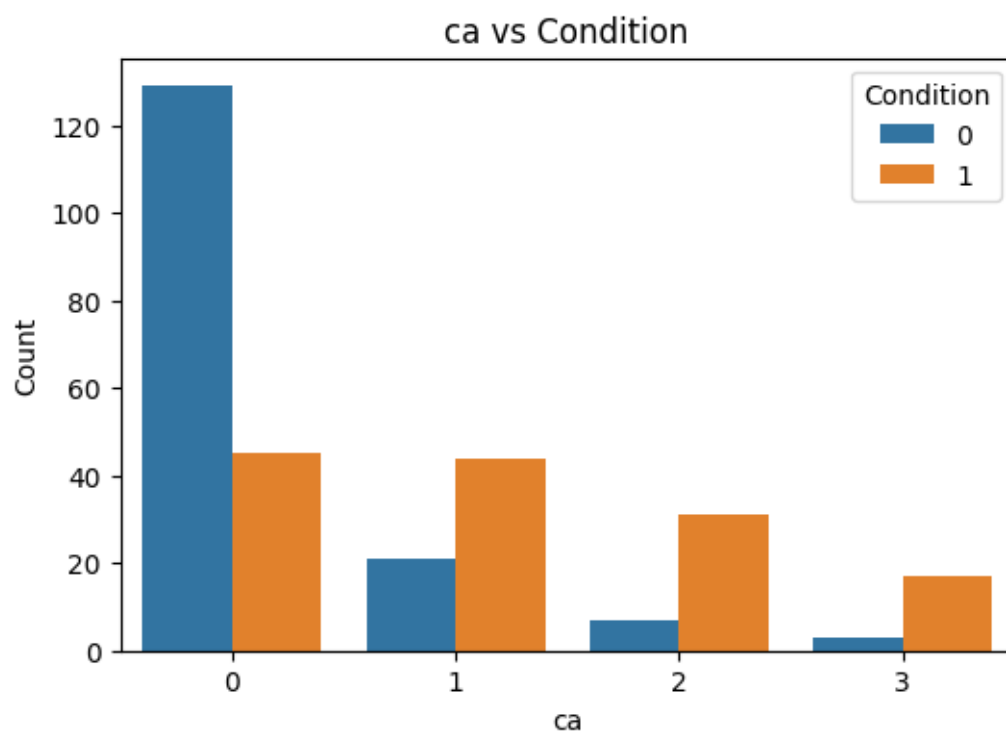
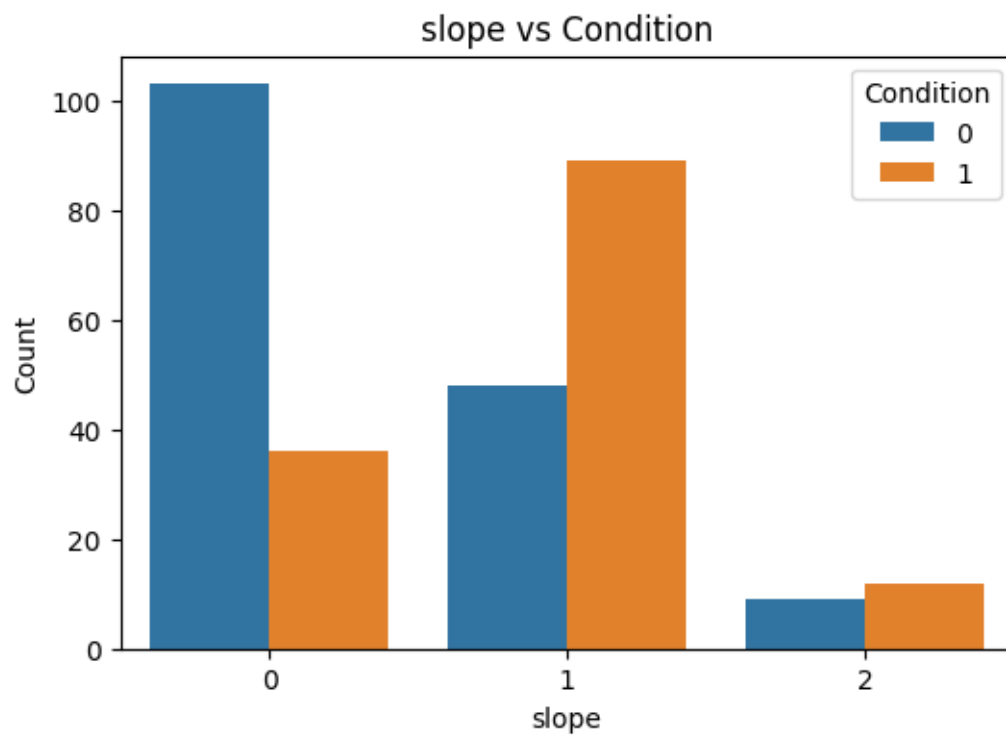
```
[60]: categorical_cols = ['sex', 'cp', 'fbs', 'restecg', 'exang', 'slope', 'ca', 'thal']

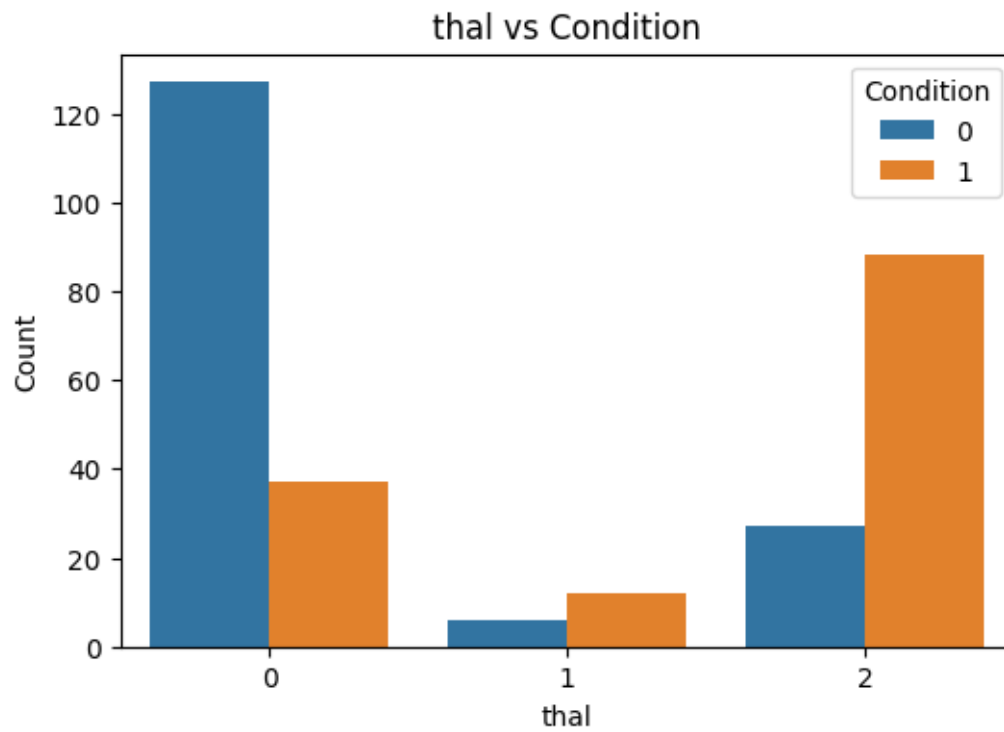
for col in categorical_cols:
    plt.figure(figsize=(6, 4))
    sns.countplot(x=col, hue='condition', data=data)
    plt.title(f'{col} vs Condition')
    plt.xlabel(col)
    plt.ylabel('Count')
    plt.legend(title='Condition')
    plt.show()
```












15 Correlation Heatmap

```
[49]: plt.figure(figsize=(10,8))
sns.heatmap(data.corr(), annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Feature Correlation Matrix')
plt.show()
```



16 Prepare Data for Modeling

```
[52]: X = data.drop('condition', axis=1)
      y = data['condition']

      # Train-test split (80% train, 20% test)
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪ random_state=42, stratify=y)
```

17 Train Logistic Regression Model

```
[53]: model = LogisticRegression(max_iter=1000)
      model.fit(X_train, y_train)
```

```
[53]: LogisticRegression(max_iter=1000)
```

18 Make Predictions and Evaluate

```
[54]: y_pred = model.predict(X_test)
      y_prob = model.predict_proba(X_test)[:, 1]

      accuracy = accuracy_score(y_test, y_pred)
      roc_auc = roc_auc_score(y_test, y_prob)


      print(f'Accuracy: {accuracy:.4f}')
      print(f'ROC AUC: {roc_auc:.4f}')
```

Accuracy: 0.9167

ROC AUC: 0.9509

19 Plot Confusion Matrix


```
[55]: cm = confusion_matrix(y_test, y_pred)
      disp = ConfusionMatrixDisplay(confusion_matrix=cm)
      disp.plot(cmap='Blues')
      plt.title('Confusion Matrix')
      plt.show()
```



20 Plot ROC Curve

```
[56]: fpr, tpr, thresholds = roc_curve(y_test, y_prob)

plt.plot(fpr, tpr, label=f'AUC = {roc_auc:.2f}')
plt.plot([0,1], [0,1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.show()
```



21 Feature Importance (Logistic Regression Coefficients)

```
[57]: features = X.columns
      coefficients = model.coef_[0]

      importance_df = pd.DataFrame({'Feature': features, 'Coefficient': coefficients})
      importance_df['AbsCoefficient'] = importance_df['Coefficient'].abs()
      importance_df = importance_df.sort_values(by='AbsCoefficient', ascending=False)

      print("Important Features (based on logistic regression coefficients):")
      print(importance_df[['Feature', 'Coefficient']])
```


Important Features (based on logistic regression coefficients):

	Feature	Coefficient
11	ca	1.046757
8	exang	0.812330
1	sex	0.675442
12	thal	0.623016
5	fbs	-0.582774

2	cp	0.498075
9	oldpeak	0.379411
10	slope	0.320660
6	restecg	0.139551
7	thalach	-0.023170
3	trestbps	0.017056
0	age	-0.013472
4	chol	0.002924

22 Visualize Feature Importance

```
[58]: plt.figure(figsize=(10,6))
sns.barplot(x='Coefficient', y='Feature', data=importance_df)
plt.title('Feature Importance - Logistic Regression Coefficients')
plt.show()
```



```
[ ]:
```

```
[62]: import getpass
username = "zoya4477"
token = getpass.getpass("Enter your GitHub token: ")
```

Enter your GitHub token:

```
[63]: !git add .  
      !git commit -m "AI-ML Tasks"
```

```
fatal: not a git repository (or any of the parent directories): .git  
fatal: not a git repository (or any of the parent directories): .git
```

```
[64]: !git pull https://github.com/zoya4477/AI-ML.git  
      !git push https://{username}:{token}@github.com/zoya4477/AI-ML.git
```

```
fatal: not a git repository (or any of the parent directories): .git  
fatal: not a git repository (or any of the parent directories): .git
```

```
[ ]:
```